

■ ForgeBench

AI Agent Harness & Evaluation Laboratory

에이전트의 성공뿐 아니라 실패한 주장까지 재현 가능하게 만든다

Codex 기반 planning · tool loop · recovery · risk routing · replay · hidden evaluation firewall

150

hash-bound evaluation targets


120/120

model-policy replays sealed

130

automated regression tests

SYSTEM



동일 base workspace · frozen model/config · SHA-256 replay binding · private grader outside Git

HELD-OUT RESULT / PRIMARY CLAIM REJECTED

Policy	Hidden success	Model	Input tokens	Safety
Accept-All	0 / 30	0	0	0
Verify-All	0 / 30	30	1,690,849	1
Random-k	0 / 30	9	348,434	0
Probe-All	0 / 30	0	0	0
Risk-Direct	0 / 30	9	474,452	0
Risk-Hierarchical	0 / 30	9	504,339	0

P5: 9 model calls · 504,339 input tokens · 0 recoveries

Risk recall 0.30 · task-cluster bootstrap 95% CI [0.0, 0.6]

WHAT THE FAILURE TAUGHT

- 모든 base가 false completion: single-class 한계
- P2/P4/P5가 같은 9개 base 선택: allocation 붕괴
- 9개 probe 모두 unsupported: short-circuit 0회
- test-time compute만 늘려서는 repair가 되지 않음

AI AGENT ENGINEER EVIDENCE

- Planning / Loop: plan-completion-repair contract
- Tool / Safety: isolated workspace-protected evidence
- Evaluation: shared base-hidden firewall-bootstrap
- Harness R&D: routing degeneracy와 실패 원인 분리

핵심 주장: 정책이 성공했다가 아니라, 성공하지 않았음을 데이터 누수 없이 증명하는 평가 시스템을 만들었다.