

■ ForgeBench

AI Agent Harness & Evaluation Laboratory

에이전트의 성공뿐 아니라 실패한 주장까지 재현 가능하게 만든다

Codex 기반 planning · tool loop · recovery · risk routing · replay · hidden evaluation firewall

90

frozen V2 Codex slots


84

externally graded V2 bases

207

automated regression tests

SYSTEM



frozen model/config · SHA-256 catalog/corpus/oracle · private grader and labels outside Git

FINAL V2 SCREENING / PREREGISTERED GATE NOT MET

Measure	Wave 4	Combined	Required	Decision
Frozen slots	30	90	-	complete
Eligible bases	28	84	-	sealed
False completions	0	10	-	preserved
False-completion clusters	0	4	6	insufficient
Failure mechanisms	0	3	4	insufficient
Passing-control clusters	10	27	6	passed

Stage B blocked · Wave 5 prohibited · thresholds unchanged

Three identical external evaluations · 0 infrastructure failures

WHAT THE FAILURE TAUGHT

- V1: 추가 compute로 ineffective repair를 해결 못함
- V2: target failure 다양성이 먼저 필요함
- raw failure 수는 cluster/mechanism을 대체 못함
- terminal rule이 outcome-driven 확장을 차단함

AI AGENT ENGINEER EVIDENCE

- Planning / Loop: plan-completion-repair contract
- Tool / Safety: isolated workspace-protected evidence
- Evaluation: 90 slots-3회 외부 평가-hash seal
- Benchmark: 30 lineages-public/private firewall

핵심 주장: 성공 점수가 아니라, 지원되지 않는 연구를 데이터 누수 없이 중단할 수 있는 평가 시스템을 만들었다.